

---

# Physical optimization (Intermediate results, Join)

# Knowledge Objectives

---

1. Explain and exemplify the two main assumption that most DBMSs do on estimating the cardinality of query results
2. Explain the need and problem of gathering statistics
3. Enumerate five join algorithms
4. Explain the prerequisites of each join algorithm (i.e. Row Nested Loops, Block Nested Loops, Sort-Match, Hash Join)
5. Write the pseudo-code of five join algorithms

# Understanding Objectives

---

1. Calculate the approximate size of a table, given its number of tuples, its structure, and the number of tuples and entries that fit in a disk block
2. Estimate the number of tuples in the output of a query, given the schema of the database and the tuples in the tables
3. Calculate the number of blocks necessary to store a given number of tuples, knowing the size of the attributes of each tuple and the bytes of each block
4. Identify when a join algorithm can or cannot be used in a given query or process tree
5. Given the statistics of the tables, estimate the cost of a join using each of the algorithms

# Physical optimization

---

Consists of **generating** the **execution plan** of a query (from the best syntactic tree) considering:

- ❑ Physical structures
- ❑ Access paths
- ❑ Algorithms

# Process tree

---

This is the tree associated to the syntactic tree that models the execution strategy

## □ Nodes

- Root: Result
- Leaves: Tables (or Indexes)
- Internal: Intermediate tables generated by a physical operation

## □ Edges

- Denote direct usage

# Physical operations

---

- ❑ Related to relational algebra
  - Physical selection: Selection [+ projection]
  - Physical join: Join [+ projection]
  - Set operations:
    - ❑ Union [+ projection]
    - ❑ Difference [+ projection]
- ❑ Other operations:
  - Duplicate removal
  - Sorting
  - Grouping and calculating aggregates

# Cost-based optimization

---

- ❑ The cost of the process tree is the sum of costs of each physical operation
- ❑ The cost of each operation is the sum of
  - Cost of solving it
  - Cost of writing its result
- ❑ Cost factors:
  - CPU
  - Memory access time
  - Disk access time

# Phases to find the minimum cost tree

---

- ❑ **Phase 1:** Alternatives generation
- ❑ **Phase 2:** Intermediate results cardinality and size estimation
- ❑ **Phase 3:** Cost estimation for each algorithm and access path
- ❑ **Phase 4:** Choose the best option and generate the access plan



# Phase 1 Execution alternatives


---

- ❑ Join order (combinatorial explosion)
- ❑ Access path to each tuple (depending on available structures)
- ❑ Execution algorithm for each operation
- ❑ Materialization or not of intermediate results (we will assume that they are always materialized)

# Join order


- We can generate different process trees by using the associative property of joins

PT 1



$(W*V)*P$

PT 2




$W*(V*P)$

## Phase 2 Cardinalities estimation

- It is based on the selectivity factor ( $0 \leq SF \leq 1$ )
  - Next to 0 means very selective (Eg: ID)
  - Next to 1 means little selective (Eg: Sex)
- SF is only needed for selection and join
  - Estimated cardinality for selection  
 $|\text{selection}(R)| = SF * |R|$
  - Estimated cardinality for join:  
 $|\text{join}(R,S)| = SF * |R| * |S|$
  - Estimated cardinality for union:
    - With repetitions:  
 $|\text{union}(R,S)| = |R| + |S|$
    - Without repetitions:  
 $|\text{union}(R,S)| = |R| + |S| - |\text{join}(R,S)|$
  - Estimated cardinality for difference (anti-join):  
 $|\text{difference}(R,S)| = |R| - |\text{join}(R,S)|$
- It is calculated from leaves to root
- Statistics are needed to estimate cardinalities

$$|R'| = |R| * SF$$



# Statistics

---

- ❑ DBA is responsible for the statistics to be fresh
- ❑ Example of kinds of statistics
  - Regarding relations:
    - ❑ Cardinality
    - ❑ Number of blocks
    - ❑ Average length of records
  - Regarding attributes:
    - ❑ Length
    - ❑ Domain cardinality (maximum number of different values)
    - ❑ Number of existing different values
    - ❑ Maximum value
    - ❑ Minimum value
- ❑ Main hypothesis in most DBMS
  - Uniform distribution of values for each attribute
  - Independence of attributes

# Statistics in Oracle 10g

---

a) `ANALYZE [TABLE|INDEX|CLUSTER] <name>  
[COMPUTE|ESTIMATE] STATISTICS;`

```
ANALYZE TABLE departments COMPUTE STATISTICS;  
ANALYZE TABLE employees COMPUTE STATISTICS;
```

b) `DBMS_STATS.GATHER_TABLE_STATS( <esquema>,  
<table> );`

```
DBMS_STATS.GATHER_TABLE_STATS("username","departments");  
DBMS_STATS.GATHER_TABLE_STATS("username","employees");
```

# Selectivity factor of a Selection (I)

- General rule: favorable cases / possible cases
- Assuming equi-probability of values
  - $SF(A=c) = 1/ndist(A)$
- Assuming uniform distribution and  $A \in [\min, \max]$ 
  - $SF(A > c) = (\max - c) / (\max - \min)$  (Not for integers)
    - $SF(A > c) = 0$  (if  $c \geq \max$ )
    - $SF(A > c) = 1$  (if  $c < \min$ )
  - $SF(A > v) = 1/2$
  - $SF(A < c) = (c - \min) / (\max - \min)$  (Not for integers)
    - $SF(A < c) = 1$  (if  $c > \max$ )
    - $SF(A < c) = 0$  (if  $c \leq \min$ )
  - $SF(A < v) = 1/2$
- Assuming  $ndist(A)$  big enough
  - $SF(A \leq x) = SF(A < x)$
  - $SF(A \geq x) = SF(A > x)$

# Selectivity factor of a Selection (II)

- Assuming P and Q statistically independent
  - $SF(P \text{ AND } Q) = SF(P) \cdot SF(Q)$
  - $SF(P \text{ OR } Q) = SF(P) + SF(Q) - SF(P) \cdot SF(Q)$
- $SF(\text{NOT } P) = 1 - SF(P)$
- $SF(A \text{ IN } (c_1, c_2, \dots, c_n)) = \min(1, n / \text{ndist}(A))$
- $SF(A \text{ BETWEEN } c_1 \text{ AND } c_2) = \frac{\min(c_2, \text{max}) - \max(c_1, \text{min})}{(\text{max} - \text{min})}$  (Not for integers)
- $SF(A \text{ BETWEEN } v_1 \text{ AND } v_2) = 1/4$
- $SF(A \text{ BETWEEN } c_1 \text{ AND } v_2) = 1/2 SF(A > c_1)$
- $SF(A \text{ BETWEEN } v_1 \text{ AND } c_2) = 1/2 SF(A < c_2)$

# Selectivity factor of a Join (I)




---

- ❑ It is difficult to approximate the general case  $R[A \theta B]S$
- ❑ Depending on the comparison operator:
  - $SF(R[A \times B]S) = 1$  ;  $SF(R[A < > B]S) = 1$
  - $SF(R[A = B]S) = 1/\max(\text{ndist}(A), \text{ndist}(B))$ 
    - ❑ There is no FK
    - ❑ One domain is subset of the other one
  - $SF(R[A < B]S) = 1/2$



# Selectivity factor of a Join (II)

- $SF(R[A=B]S) = 1/|R|$ 
  - S.B is FK to R.A
  - S.B is not null
  - R.A is PK
  
- $SF(R'[A=B]S) = (1/|R'|) * sfr$   
 $= 1/|R|$ 
  - $R' = \text{selection}(R)$  with  $SF = sfr$
  - S.B is FK to R.A
  - S.B is not null
  - R.A is PK
  
- $SF(R[A=B]S') = (1/|R|)$ 
  - $S' = \text{selection}(S)$  with  $SF = sfr$
  - S.B is FK to R.A
  - S.B is not null
  - R.A is PK



# Intermediate results estimation

---

- ❑ Besides the cost of executing each operation, we also need to know the cost of writing intermediate results into a temporal file
  - Record length
    - ❑  $\sum \text{attribute length}_i$  (+ control information)
  - Number of records per block
    - ❑  $R_R = \lfloor \text{block size} / \text{record length} \rfloor$
  - Number of blocks per table
    - ❑  $B_R = \lceil |R| / R_R \rceil$

# Phase 3

---

- Join
- Selection
- Sort
- Projection

# Join algorithms

---

- ❑ Scan (Clustered structure)
- ❑ Row Nested Loops
- ❑ Block Nested Loops
- ❑ Sort-Match
- ❑ Hash Join

# Clustered Structure

## □ Space

- $\lceil 1.5B_{RS} \rceil$

$$B_{RS} = B_R + B_S$$

$$R_{RS} = (|R| + |S|) / (B_R + B_S)$$

## □ Access paths

- Cost of table scan of R
  - $\lceil 1.5B_{RS} \rceil$
- Cost of table scan of S
  - $\lceil 1.5B_{RS} \rceil$

Dept 1	Employee 14	Employee 8	Employee 6		Dept 2	Employee 3		Dept 3		Dept 4	Employee 18	Employee 2	
--------	-------------	------------	------------	--	--------	------------	--	--------	--	--------	-------------	------------	--

# Row Nested Loops

## □ Algorithm

```

for each block of R
  read block of R into a memory page
  for each tuple t in the read page
    if there is an index in S for attribute A
      go through the index of S.A using the value of t.A
      if there is any tuple satisfying the join condition
        if we are interested in attributes of S
          go to the corresponding tuples of S
        endIf
        generate result
      endIf
    else scan the whole table S and generate result
  endIf
endForEach
endForEach

```

$u = \%load \cdot 2d = (2/3) \cdot 2d$   
 $h = \lceil \log_u |T| \rceil - 1$   
 $k = \text{average appearance of each value of R.A in S.A}$

## □ Cost, if we do NOT look for attributes of S (semi-join)

- B+:  $B_R + |R| \cdot (h_S + (k-1)/u_S)$
- Hash:  $B_R + |R|$

## □ Cost, if we DO look for attributes of S

- B+:  $B_R + |R| \cdot (h_S + (k-1)/u_S + k)$
- Clustered:  $B_R + |R| \cdot (h_S + 1 + 1.5(k-1)/R_S)$
- Hash:  $B_R + |R| \cdot (1 + k)$

## □ Considerations

- It is only useful if there is an index over the join attribute
- A hash index can only be used for equi-join
- We cannot use this algorithm if we performed previous operation over the table
- $B_R$  is the real number of blocks (taking into account possible cluster)

# Block Nested Loops (I)

## □ Algorithm

repeat

  read M blocks of R

for each block of S

    read block of S into a memory page

for each tuple t in the pages of R

for each tuple s in the page of S

if (t.A  $\theta$  s.B) then generate result

endIf

endForEach

endForEach

endForEach

until no more blocks to read from R

## □ Cost (with M+2 memory pages)

- $B_R + B_S \cdot \lceil B_R/M \rceil$


## □ Considerations

- It can always be used (even for  $\theta$ -join)
- It is of special interest if  $B_R \leq M$
- It is not symmetric
  - It is better if the smaller table is in the external loop
  - $B_R$  and  $B_S$  are the real number of blocks (taking into account possible cluster)


# Block Nested Loops (II)

$M=2$

Accesses



$R_1 R_2 R_3 R_4$   
 $S_1 S_2$



$R_1 R_2$   
 $S_1 S_2 S_3 S_4$



# Block Nested Loops (II)

M=2

E	E	I	O
$R_1$	$R_2$	$S_1$	

Accesses

R  
R  
R

E	E	I	O

$R_1 R_2 R_3 R_4$   
 $S_1 S_2$

$R_1 R_2$   
 $S_1 S_2 S_3 S_4$

# Block Nested Loops (II)

M=2

Accesses

R  
R  
R

E	E	I	O
$R_1$	$R_2$		

$R_1 R_2 R_3 R_4$   
 $S_1 S_2$

E	E	I	O

$R_1 R_2$   
 $S_1 S_2 S_3 S_4$

# Block Nested Loops (II)

M=2

Accesses

R  
R  
R  
R


E	E	I	O
R <sub>1</sub>	R <sub>2</sub>	S <sub>2</sub>	

R<sub>1</sub> R<sub>2</sub> R<sub>3</sub> R<sub>4</sub>  
S<sub>1</sub> S<sub>2</sub>

E	E	I	O

R<sub>1</sub> R<sub>2</sub>  
S<sub>1</sub> S<sub>2</sub> S<sub>3</sub> S<sub>4</sub>

# Block Nested Loops (II)



# Block Nested Loops (II)

M=2

Accesses

E	E	I	O
$R_3$	$R_4$	$S_1$	


$R_1 R_2 R_3 R_4$   
 $S_1 S_2$

R  
R  
R  
R  
R  
R  
R

E	E	I	O

$R_1 R_2$   
 $S_1 S_2 S_3 S_4$

# Block Nested Loops (II)



# Block Nested Loops (II)

M=2

E	E	I	O
$R_3$	$R_4$	$S_2$	

$R_1 R_2 R_3 R_4$   
 $S_1 S_2$


Accesses

R  
R  
R  
R  
R  
R  
R  
R

E	E	I	O


$R_1 R_2$   
 $S_1 S_2 S_3 S_4$

# Block Nested Loops (II)







# Block Nested Loops (II)




# Block Nested Loops (II)




# Block Nested Loops (II)



# Block Nested Loops (II)



# Block Nested Loops (II)



# Sort-Match (I)

## □ Algorithm (for unique join attributes)

```

Sort R and S (if necessary)
 $t_R := \text{first}(R)$ ;  $t_S := \text{first}(S)$ ;
while not (end(R) or end(S))
  if ( $t_R[A] < t_S[A]$ )  $t_R := \text{next}(R)$ ;
  elsif ( $t_R[A] > t_S[A]$ )  $t_S := \text{next}(S)$ ;
  else generate result from  $t_R$  and  $t_S$ ;  $t_R := \text{next}(R)$ ;  $t_S := \text{next}(S)$ ;
endif
endWhile

```

## □ Cost, with $M+1$ memory pages


- Sorting R (same for S):
  - If sorted: 0
  - Elsif  $B_R \leq M$ :  $2B_R$
  - Else:  $2B_R \cdot \lceil \log_M B_R \rceil$
- Merging R and S:  $B_R + B_S$

## □ Considerations

- It is only useful for equi-join,  $<$ -join,  $>$ -join and anti-join
  - The given cost corresponds to equi-join and anti-join
- It is of special interest when at least one table is Clustered
  - If both tables are sorted, we only need 3 memory pages
- The result is already sorted
- $B_R$  and  $B_S$  are the real number of blocks (taking into account possible cluster). Beware: after sorting, the auxiliary tables have no empty space.

# Sort-Match (II)

Accesses



1	2
2	4
3	5
5	6

# Sort-Match (II)

Accesses

R  
R

1 2	2 4	
--------	--------	--



3 5	5 6
--------	--------



# Sort-Match (II)

Accesses



R  
R



# Sort-Match (II)

Accesses

R  
R





# Sort-Match (II)

---

Accesses

R  
R  
R





# Sort-Match (II)

---

Accesses

R  
R  
R





# Sort-Match (II)

---

Accesses

R  
R  
R



# Sort-Match (II)

---

Accesses

R  
R  
R  
R


5	5 6	2
---	--------	---

# Sort-Match (II)

---

Accesses


R  
R  
R  
R



# Sort-Match (II)

Accesses

R  
R  
R  
R



2  
5




# Sort-Match (II)

---

Accesses

R  
R  
R  
R



2  
5

# Hash Join (I)

## □ Algorithm

Partition R into  $p$  parts and redistribute the tuples using a hash function

Partition S into  $p$  parts and redistribute the tuples using the same hash function

(partitions must assure that each part of one relation fits into  $M$  memory pages)

Use Block Nested Loops  $p$  times (part by part)

## □ Considerations

- It can only be used for equi-join
- We will take  $p = \lceil B_{\text{Smaller}} / M \rceil$ . The size of each part of  $B_{\text{Smaller}}$  is  $M$  (assuming hash results in uniform distribution of values). Instead, the size of the other is bigger.
- If  $B_{\text{Smaller}} \leq M$ , the algorithm coincides with Block Nested Loops ( $p = 1$ )

## □ Cost, with $M+2$ memory pages

- If  $B_{\text{Smaller}} \leq M$ :  $B_R + B_S$
- If  $M < B_{\text{Smaller}} \leq M^2 + M$ :  $2B_R + 2B_S + B_R + B_S$  ( $p \leq (M^2 + M) / M = M + 1$ )
- If  $B_{\text{Smaller}} > M^2 + M$ :
- $B_R$  and  $B_S$  are the real number of blocks (taking into account possible cluster). Beware: after redistribution, the tables have no empty space.

# Hash Join (II)

## Accesses

$$M=1$$
$$B_1=B_2=2$$




Table 1

1
2
3
5

Table 2

2
4
5
6

# Hash Join (II)

## Accesses

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

1
2
3
5

Table 2

2
4
5
6

# Hash Join (II)

## Accesses

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O E

1  
2

3  
5

Table 2

2  
4

5  
6

# Hash Join (II)

Accesses

R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O E

3  
5

Table 2

2  
4

5  
6

# Hash Join (II)

Accesses

R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O E

3  
5

Table 2

2  
4

5  
6

# Hash Join (II)

# Accesses

R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O E

3
5

Table 2

2
4

5
6



# Hash Join (II)

Accesses

R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1  
O E

Table 2



# Hash Join (II)

# Accesses

R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1  
O E

Table 2

2
4
5
6

# Hash Join (II)

Accesses

R  
R  
W

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O E

1
3

Table 2

2
4

5
6

# Hash Join (II)

## Accesses

R  
R  
W

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O E

1
3

Table 2

2
4

5
6

# Hash Join (II)

Accesses

R  
R  
W  
W  
W

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
1	2
3	
5	

Table 2

O	E
	2
	4
	5
	6

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O   E

1	2
3	
5	

Table 2

O   E

5	
6	

# Hash Join (II)

Accesses

R  
R  
W  
W  
W  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O E

1	2
3	

5
---

Table 2

O E

5
6

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O   E

1	2
3	
5	

Table 2

O   E

5	
6	



# Hash Join (II)

Accesses

R  
R  
W  
W  
W  
R  
W

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
1	2
3	
5	

Table 2

O	E
	2
	4
5	6

# Hash Join (II)

Accesses

R  
R  
W  
W  
W  
R  
W  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
1	2
3	
5	

Table 2

O	E
	2
	4

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
1	2
3	
5	

Table 2

O	E
	2
	4

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
1	2
3	
5	

Table 2

O	E
	2
	4

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
R  
W  
W

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
1	2
3	
5	

Table 2

O	E
5	2
	4
	6

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
R  
W  
W  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
	2
5	

Table 2

O	E
	2 4
6	

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
R  
W  
W  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
	2
5	

Table 2

O	E
	2
	4
6	

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
R  
W  
W  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
	2
5	

Table 2

O	E
	2
	4
6	



# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
R  
W  
W  
R  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
	2

Table 2

O	E
	2
	4
	6

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
R  
W  
W  
R  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1

O	E
	2

Table 2

O	E
	2
	4
	6

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
W  
R  
W  
W  
R  
R  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1  
O E

Table 2  
O E

6

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
W  
R  
W  
W  
R  
R  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$




Table 1  
O E

Table 2  
O E


6

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
W  
R  
W  
W  
R  
R  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$



5
2

Table 1  
O E

Table 2  
O E


6
---

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
W  
R  
W  
W  
R  
R  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$



5  
2

Table 1  
O E

Table 2  
O E


6

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
W  
R  
W  
W  
R  
R  
R  
R  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$



5
2

Table 1  
O E


Table 2  
O E

# Hash Join (II)

## Accesses

R  
R  
W  
W  
W  
R  
W  
R  
W  
W  
R  
R  
R  
R  
R  
R

$M=1$   
 $B_1=B_2=2$   
 $p=2$



5  
2

Table 1  
O E

Table 2  
O E



# Summary table

---

	No index	B+	Hash	Clustered	Clustered structure
All tuples	Scan				
One tuple		Go through index Go to table	Apply function Go to bucket Go to table	Go through index Go to table	
Several tuples		Go through index Follow leaves Go to table		Go through index Go to table Scan table	
Join	Block Nested Loops Or Hash Join	Row Nested Loops	Row Nested Loops	Row Nested Loops Or Sort-Match	Scan

# Summary

---

- ❑ Cost-based optimization
  - Alternatives in the structures
  - Alternatives in the execution
    - ❑ Selection algorithms
    - ❑ External sort algorithms
    - ❑ Projection algorithms
    - ❑ Join algorithms
  - Intermediate results estimation

# Bibliography

---

- G. Gardarin and P. Valduriez. *Relational Databases and Knowledge Bases*. Addison-Wesley, 1998
- R. Ramakrishnan and J. Gehrke. *Database Management Systems*. McGraw-Hill, 3<sup>rd</sup> Edition, 2003
- S. Lightstone, T. Teorey and T. Nadeau. *Physical Database Design*. Morgan Kaufmann, 2007
- J. Sistac. *Sistemes de Gestió de Bases de Dades*. Editorial UOC, 2002
- Y. Ioannidis. *Query Optimization*. ACM Computing Surveys, vol. 28, num. 1, March 1996.
- J. Lewis. *Cost-Based Oracle Fundamentals*. Apress, 2006